

Implementation details of core.async Channels

- Speaker: Rich Hickey
- Conference: [EuroClojure](#) - June 2014
- Video: <https://vimeo.com/100518968>



Inside core.async Channels

Rich Hickey

Figure 1: 00.00.00 Inside

All right; now to the task at hand.

So, what I want to talk about today is sort of some of the implementation details of core.async channels, mostly just so that more people know about them and how it works. And, as a point of interest, and I was looking forward to giving a technical talk as opposed to trying to be inspirational.

So, implementation details - oh, there's something else I want to say about the prior talk. Everything about that may be class loader or whatever. In the fast load branch, there's now a cache in front of that, and that is no longer a hotspot for the compiler. So if you just look in that branch, you'll see that it's already been improved and that will be in the next version of Clojure. There's a bunch of other optimizations in the fast load branch, which will be interesting to people who want Clojure to load faster.

So there are going to be implementation details in this talk, so this is not something you want to marry or, you know, really get fixated upon. All of it is subject to change. We're still not really advertising some of the protocols I'm going to be talking about as extension points, but it's useful to understand how they work for you as a user of the library.

Warning!

- Implementation details
- Subject to change

Figure 2: 00.00.04 Warning!

The Problems

- Single channel implementation
- For use from both dedicated threads and go threads
simultaneously, on same channel
- alt and atomicity
- multi-read/write
- concurrency

Figure 3: 00.01.13 The Problems

So these are the problems we're trying to solve, or I was trying to solve, in taking on channels. The idea was to support both genuine threads and lightweight threads, but I knew a couple things that I wanted to do. One was support a single implementation of a channel. That was to say, there is one channel and, whether you were consuming or producing for it with genuine threads or with the go, lightweight go threads should be independent. We don't care. And it should be able to support readers and writers from either side, so there's a strict orthogonality between the channel itself and the consumers and producers and the thread stuff that's in core.async, which I'm not going to talk about too much.

So we want to be able to use from both dedicated threads and go threads simultaneously with the same channel. The other thing I wanted to support was alt with full atomicity. If you'd ever looked at some of the Java CSP libraries, alt was something that was sort of not really that great. And so we wanted to do that, and that's a little bit tricky to do atomically.

We want to support multireader and multi-writer and make sure all the concurrency is correct. And, like most of the other Clojure constructs, we make it so that the construct deals with all the ick of threads and mutexes and things like that, and that you don't have to. Part of the value proposition is that that stuff is sort of in there already. You know, I write it so you don't have to.

Of course, I am going to be focusing on the Java channel implementation, the JavaScript / the ClojureScript one is similar, but it doesn't actually have a lot of the concurrency challenges because there's only one thread in the JavaScript engine.

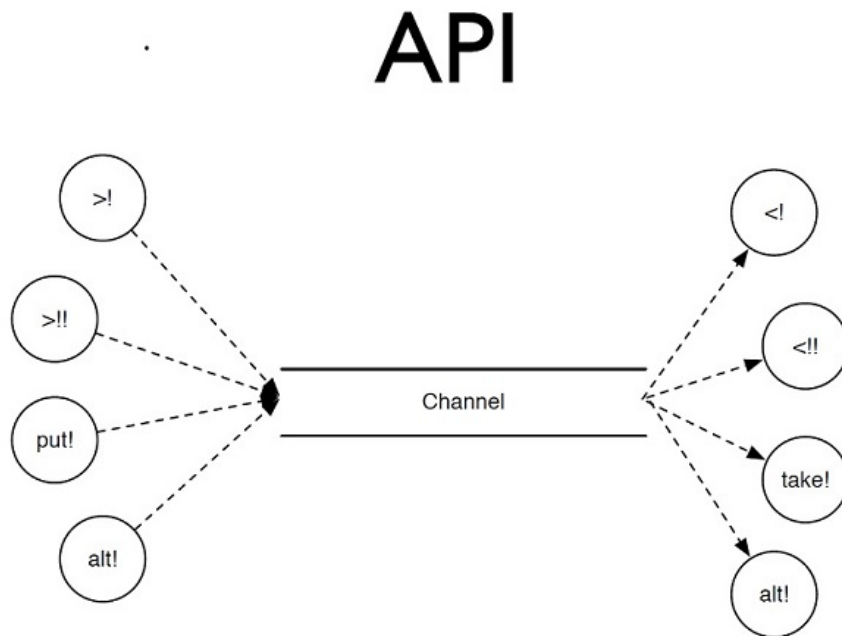


Figure 4: 00.03.02 API

So, how many people are using core.async? Yeah. Love it. And how many people have heard about it? Right. I'm not actually going to describe core.async at the API level at all today, so I apologize if some of

this surface stuff is not familiar to you yet. But the basic idea is that there are these channels, and you can put stuff in one end of the channel and you can get stuff out of the other end.

The fundamental operations are this sort of lightweight parking input, a blocking input, which blocks a real thread, a nonthreaded put. This is sort of their external entry point to a channel if you're writing a callback handler. Put bang (put!) is the right way to get data in. You don't need to start up a thread or a go block. In fact, you shouldn't. You should just use put and alt, which is sort of the underpinnings of the multi-operation thing.

Alt is really the great feature of modern CSP implementations where you can have more than one operation that you initiate, and what you will know is that, at most, one of those operations will complete, and none of the other operations will do anything. And this is sort of like an algorithmic select, if you will. You know, it's the equivalent of the sort of Unix, you know, socket capabilities.

And you have alt on both ends. And on the other end, you have the same thing. You take stuff out of the channel. It's not an RPC mechanism. It's just a conveyor belt. Put stuff on one end of the conveyor belt; you take stuff off the other end. So that's the application programmers interface through a channel. But like most things, there's usually another level of interface below.

And, you know, if you haven't heard me or someone talk about using protocols, that protocols are for service provider interfaces, that's what that's about, right? So a lot of the functions up here, they're ordinary functions. And the functions down here are functions on a protocol. And it's the contract between those API functions and any implementation of a channel. It happens to be the case that there's one really first class, full implementation of channel and some other lightweight stub-like things to help with some of the mixers and some of the other mults and things that we have in the API. But the multireader, multi-writer channel is what I'm going to be talking about today.

So the service provider interface basically boils down all of the things you can do with either end of the channel to two things: the impl/put! and impl/take! Impl/put! takes the value you want to put and handler. And, for now, you can think about that handler as just a callback. And impl/take!, similarly, returns something from the thing and it takes as an argument, again, a callback.

And I'll explain the insides of handlers in a minute. And you can define all of those operations in terms of one of those two or more than one of those operations. So alt can put and take in the same alt - you can get some stuff, try to get some stuff, and try to put some stuff and see what happens.

So, what is inside a channel? Well, when you use the channel API, you create a channel. You're able to say whether or not it takes a buffer, and you can supply an "and" and a buffer will be created, or you can actually hand in a buffer, and there are a few flavors of buffers. There are fixed buffers, and there are windowed buffers, which I'll talk about later. So inside the channel, obviously, it's got to keep track of that buffer you handed it, and that's the thing in the middle. And the thing that's in the buffer is always data. It's what was put on the channel: data values.

But if you think about the fact that the way channels work is with synchrony, right, so by default there is no buffer and a reader has to wait for a writer or a writer has to wait for a reader, and then there's this synchronous flow of control, that means that they have to wait somewhere. And, in particular, when you have multiple readers and multiple writers, they have to wait somewhere. So it ends up that there are, inside any channel, two queues, if you will, of pending operations. This is not just about the data. It's about the operation not yet completed. And they're just linkless, but they're internal to the channel and managed by it. You can't really touch them. But they will accumulate operations in the case that no one is available or there's not buffering available.

It's very important that you distinguish these lists of operations from the buffer of data. And I hope, in this talk, if you take away nothing else, that you get a good grip on how that works with channels.

So there are a bunch of things that have to remain true in implementing this and getting it right. And some are just, they just fall out of the logic of it, right? What we said was somebody trying to take the past away for somebody to put or vice versa. That means that there should never be a situation in which there are both pending puts and takes because they would have found each other, right? You can't have both people

SPI

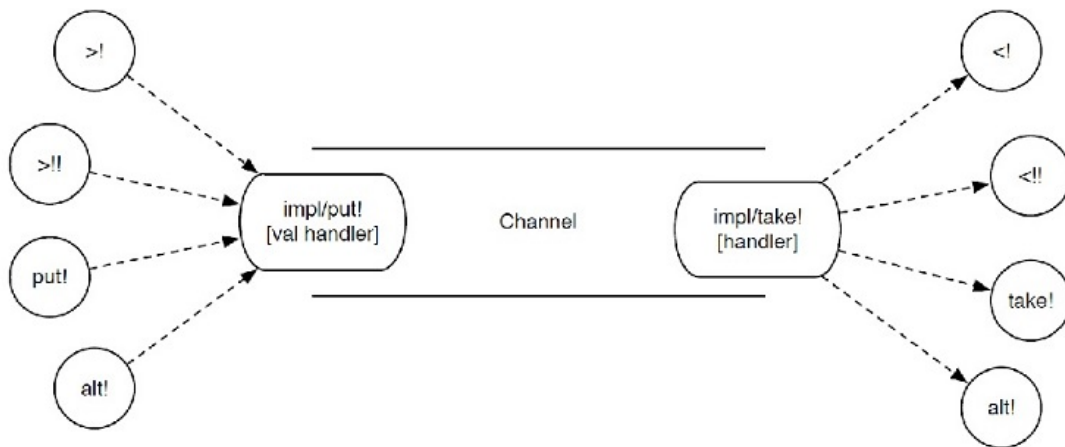


Figure 5: 00.05.37 SPI

Anatomy

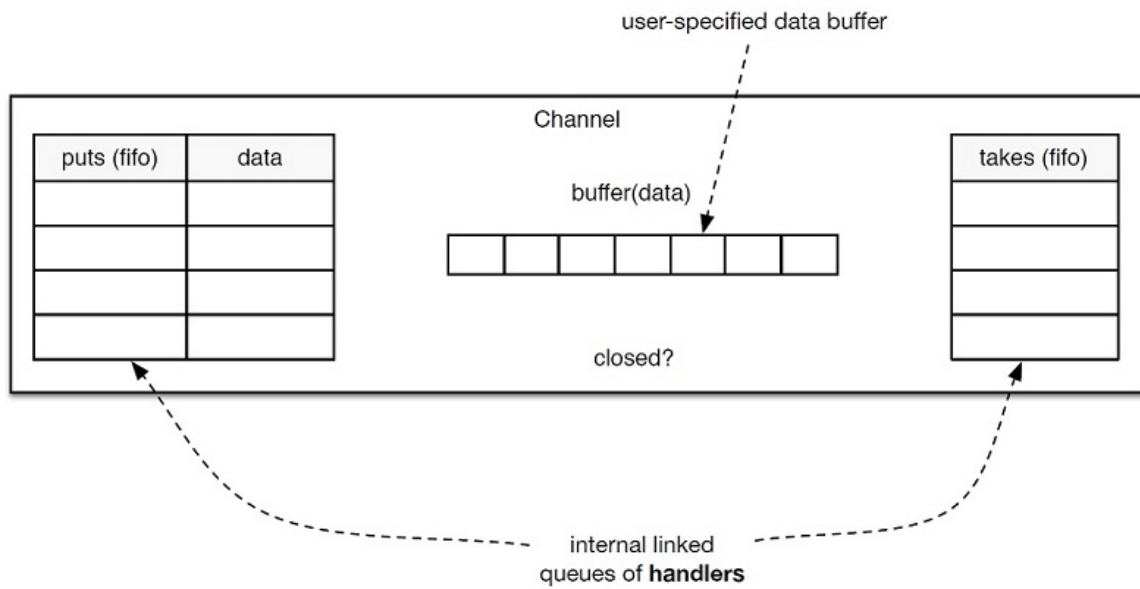


Figure 6: 00.06.23 Anatomy

Invariants

- Never pending puts and takes
- Never takes and anything in buffer
- Never puts and room in buffer
- take! and put! use channel mutex
- no global mutex

Figure 7: 00.08.03 Invariants

waiting to put and people waiting to take. If it's working, they should have met. They should have agreed to do a transfer.

You should never have pending takes if there's anything in the buffer, right, because why is it in the buffer then. Somebody came in to take, and you queued them. You should have given them something from the buffer, so that should never happen.

There should never be puts that wait in that queue if there's room in the buffer. If there's a buffer and there's room in it, a put should go and put something in the buffer and finish, not go into the wait, the waiting room.

Take and put is not an invariant except it's part of the design. Take and put use a channel wide mutex to work. And you'll see that there are actually very few conflicting states.

The other thing is, we don't want any global mutexes. At no point should there be some core.async wide mutex or even multichannel mutual exclusion, and I'll show you a little bit more about that later.

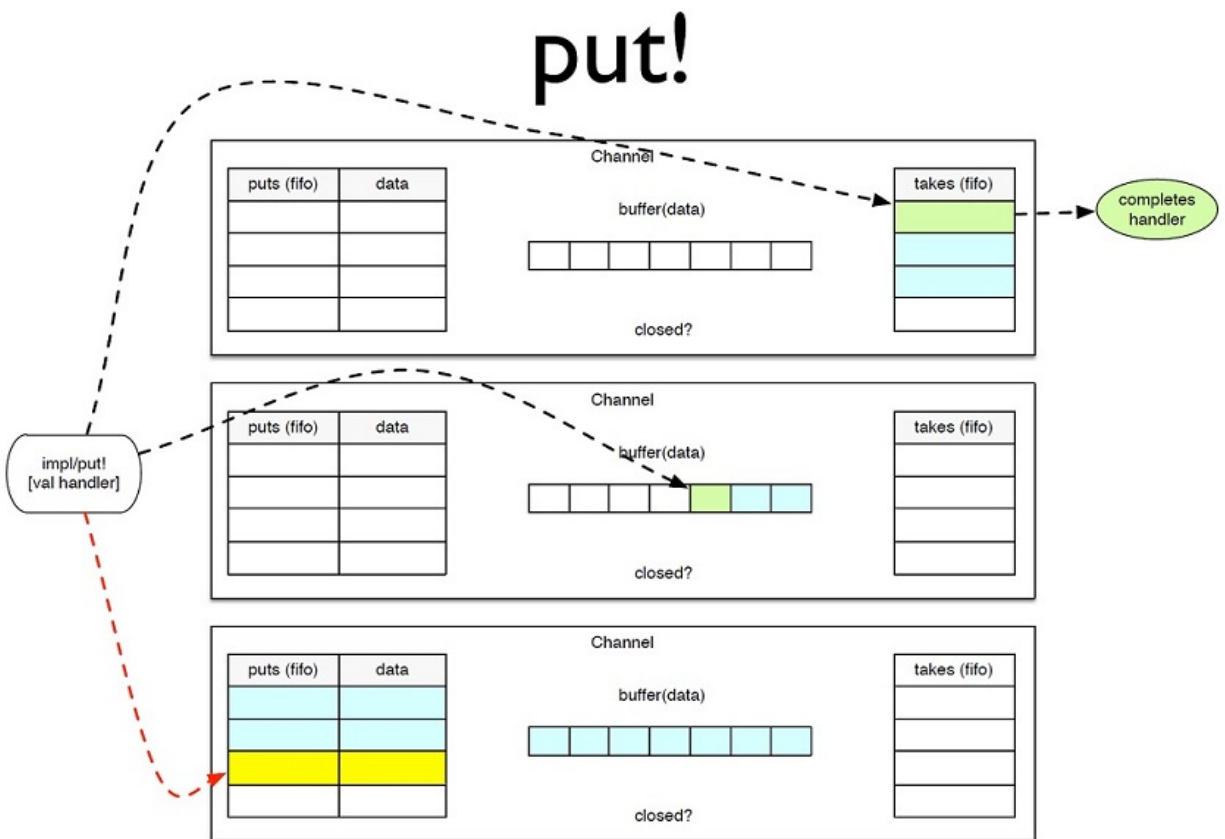


Figure 8: 00.09.33 put!

So what can happen? What are the scenarios? If we take this internal model and look at the scenarios for put, a few things could happen. In the first case, we have a put coming in, and it encounters a channel where there are one or more waiting take operations. At that point it immediately completes that waiting take, the first one in the queue. It says, okay, you're waiting for data. I have data. We're both good to go. And the put returns immediately, and the take's handler gets completed. And I'll explain handlers and completion in a little bit. So that's that scenario.

The next scenario is, well, there's stuff in the buffer. If there's stuff in the buffer and there's room in the buffer, we know there can be no pending takes and puts due to the invariants I just described. So the put is going to come in and say, room in the buffer. I'll put my stuff there and I'm out of here. I'm finished. So in that case as well, put succeeds and sort of immediately completes.

And the last scenario is that either there's no buffer or the buffer is full. But you can describe both those cases as being there's no room in the buffer. So a putter comes in, got us some new information. There's nowhere for it to go. In that case, the put gets in queue. The put operation gets in queue.

So you'll see that there are two columns in this queue because what we have to hang onto is both the data that we want to eventually convey, which we have not yet conveyed - we actually really haven't succeeded in putting it in the channel - and the handler that we need to communicate with when and if that ever gets into the channel. So there are two things always on the left-hand side. And you'll see the red arrow here implies something I'll talk more about later, which is this is actually the only case of these scenarios where we are going to have that parking effect or blocking effect because we can't finish this job right now. There's nowhere to put this stuff.

put! - windowed buffers

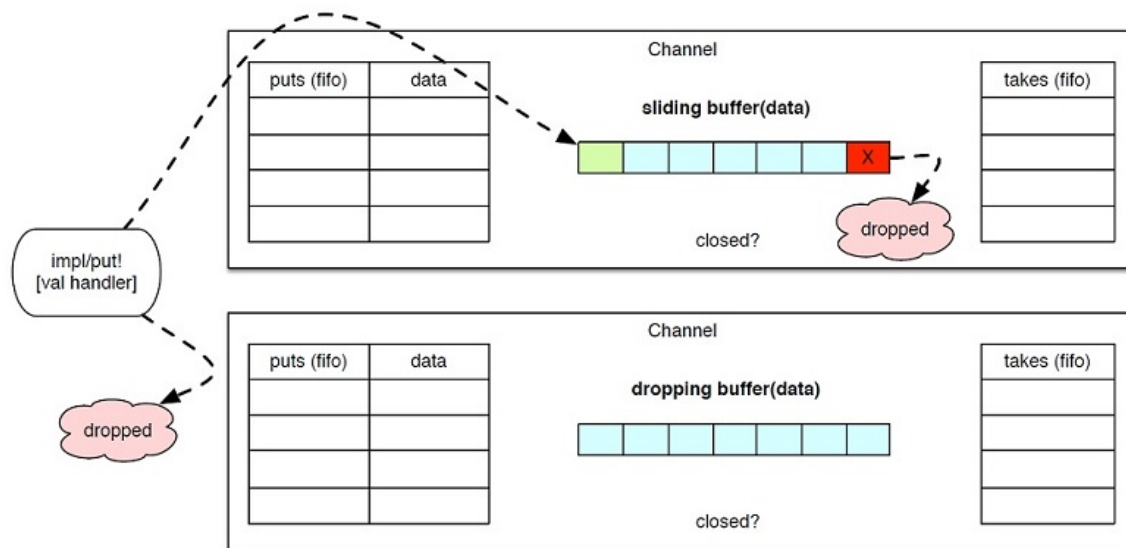


Figure 9: 00.11.48 put! - windowed buffers

I talked about the fact that there are these other buffers that are available. They're both, you know, kinds of windowing buffers. Basically, they can't keep everything, but their purpose in life is to make sure that you've installed a policy such that, especially if you're getting raw data from some external source, and you are using put bang (`put!`), there's no backpressure on put bang (`put!`). All those put bangs (`put!`) where do they go? Well, they just go in that queue. So it's a bad design for you to sort of not take care of. What should happen if I can't keep up?

One of the ways you can take care of not being able to keep up is through implement backpressure by using a blocking construct on the entry point. The other is just to use one of these windowed buffers, which say I know what I should do. If I can't keep up, I only want the latest information, and that's the first case. That's the sliding buffer. That says the latest information takes priority. If you come in with some new data, we're going to drop the first thing in the sliding window. So basically what happens is the first thing in the buffer, it gets dropped.

That means some piece of information that somebody successfully put in the channel gets thrown away. Their operation completed, but this policy says I have determined that I don't care if anybody ever sees that. It's more important to me to manage memory and to stay up to date than it is for every piece of information to get processed. And the way I'm going to do that is by dropping the oldest stuff. So that means there is room and the data goes into the buffer.

The other form of windowed buffer is one that says, once the buffer is full, anything new that comes in, throw that away. That's the dropping buffer in which case you just come in, and it says, okay, thanks for that data. I just dropped it on the floor, but you're good to go. Your operation has completed and you're fine.

These are both good, and I anticipate we'll have more sophisticated policy-based buffers, but a policy-based buffer is a great way to sort of say up front this is what I want to have happen in this scenario, as opposed to sort of saying later, "Wow, I wonder why this isn't working?" So that's how it ends in both those cases, so both those cases are also, from the perspective of the person putting, they complete right away.

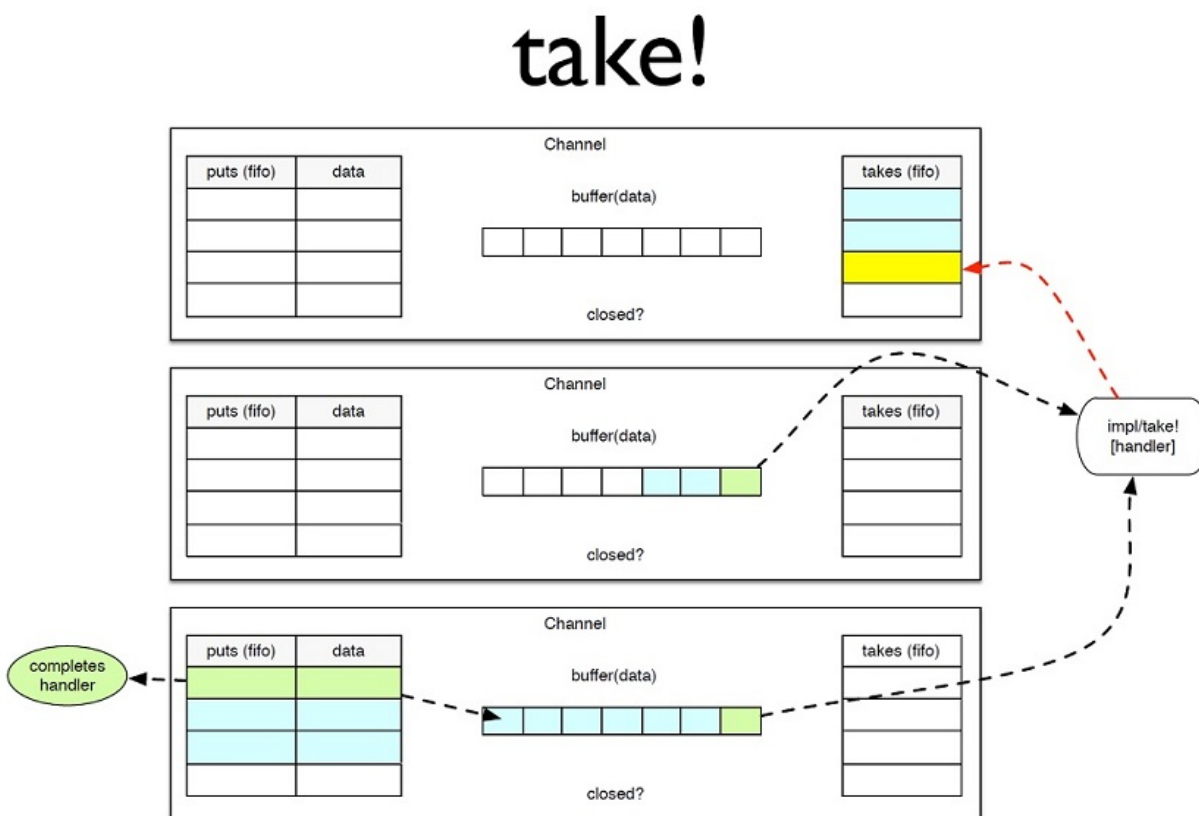


Figure 10: 00.14.06 take!

All right, and I'll take the other side. Really, it's the same three possible states of channels, right? What

happens when you take off of those three scenarios? In the first case, there are already people waiting, and there's nothing in the buffer. If I asked to take, my request has to wait. There's nothing to take. We have no stuff. There's nobody waiting to write, and there's nothing in the buffer, so your take gets in queue.

The only thing that gets in queue there is the handler. There's no data associated with the take, so that's a single column queue. Just enqueue the handler associated with the take. And that's got the red arrow. That's where that operation will park the thread that requests it.

In the next scenario, we have some data in the buffer, so we're going to hand that to the taker, and that will immediately complete. It'll get the first thing out of the buffer.

And in the third scenario, there might be no buffer or the buffer is full. In either case, and there's pending writer puts, in that case you will get something. You may get the first thing out of the buffer or, if there's no buffer, you'll get the thing from the first waiting put. But what will happen is you will get something as the taker, and the first waiting put will complete.

They may complete and hand it to you if there's no buffer, or they may complete and put it at the end of the buffer, as what's shown in this diagram. But that's another case where somebody is completing. So there's a correspondence between the red arrows in one case and the green completes handler in the other. That's sort of the magic of making this work.

close!

- all pending takes complete with nil (closed)
- subsequent puts complete with nil (already closed)
- subsequent takes consume ordinarily until empty

any pending puts complete with true

takes then complete with nil

Figure 11: 00.15.55 close!

There's the third operation, which is somewhat orthogonal to the rest, which is closing the channel. When we close a channel, all the pending takes complete right away with nil. You've got nothing. And if there

are pending takes, it means there was nothing in the buffer, and there's no pending reads due to that earlier invariant.

If you put later, you get nil. That's the already closed signifier, which is new. I don't know how long ago we added it, but it was something people wanted to know. Like, am I right? Am I wasting my time writing to this? So you can now detect that.

Subsequent takes will continue to consume data, so if there is data, if there's data inside the buffer or if there were pending puts, they'll still be delivered by a closed channel. The close is sort of after the data, if you want to think about it that way. In fact, it's after all of the operations. So that will continue to work. When they are done, then you'll start getting nils back from your takes. And all those puts that were pending, as they get consumed, they'll complete. But subsequent puts will also complete with this indicator saying already closed. All right?

Queue Limits

- puts and takes queues are not unbounded
- 1024 pending ops limit
- will throw if exceeded
- not for buffering, use buffers/windowing

Figure 12: 00.17.10 Queue Limits

So we said there was this put and take thing. That's not a recipe for unbounded buffers, right? It's a design principle of core.async that we're not going to let you have unbounded buffering. Just don't want to see it. It's just going to make for programs that don't work and bugs that show up late. So there are limits to both of those queues. They're somewhat arbitrary and it may change.

Right now it's 1,000 of pending operations. If you exceed that limit, you're just going to get an exception. And it's not a happy story. It's not like something, well, I should, you know, do. It should be something I should be seeing. If you're seeing that, you have not yet architected your system with appropriate buffering or buffering policies. You just have not said I'm set up to handle this. So that exception, if you see it, is

one that you need to examine your buffering because you should not be using these operation queues in this way.

And if you think about it, if you're really using go blocks and thread blocks and the IO operations to do these things, you're talking about having 1,000 operations in progress. Not just 1,000 pieces of data, 1,000 separate threads or activities. The place you get into this, most likely will get into this, is if you do use put bang (put!) on the edge of your system because there you haven't really thought that you were starting up a process to do a job, but these channels are using these queues thinking that they're lists of processes waiting to do work. And having more than 1,000 processes waiting to do work on a single channel means you probably are trying to do too much with too little memory assigned to the job. So use buffers and use windowing policies to make sure you don't get into this scenario, but it is something you may have to tune around.

alt(s!!)

- attempts more than one op
- on more than one channel
- without global mutex
nor multi-channel locks
- exactly one op can succeed

Figure 13: 00.19.02 alt(s!!)

All right, so alt. Alt is actually the tricky operation. Almost everything that's built into this channel implementation is there to support alt because alt is the hard part.

What alt does is it attempts more than one operation. It may attempt more than one operation and often, usually, on more than one channel. So now we have more than one operation, separate channels. We have the invariant we had from before, which we don't want any global mutex. And we don't even want any multichannel lock-in, or we don't want to say, well, if you try to alt to channels A, B, and D, we're going to lock A, B, and D for the alt because we get into all of these channel blocking intersection problems.

Also, you can't hold onto them for long enough to see the things succeed, right? Alt could block arbitrarily

long. And the other trick with alt is that exactly one operation can succeed, so you have to make sure that somehow we have a way to stop everything else from happening.

alt implications

- registration of handlers is **not** atomic
- completion might occur before registrations are finished
or any time thereafter
- completion of one alternative must ‘disable’ the others
atomically
- cleanup

Figure 14: 00.20.12 alt implications

So this implies a whole bunch of constraints on the implementation. The first is, because we’re not going to have any global mutex or even multichannel mutex, it means that the registration of handlers, so I say alt. I’m going to try to read from A, read from B, write to C. Those three operations are going to be registered, which is sort of that entry into the channel, non-atomically. That registration won’t be atomic, so I’ll have already registered something for A before I try to do anything with B. That registration is live, even as I go try to do the rest of the alt job. You cannot make the registrations atomic. That’s implied by what we saw before.

That already implies this next thing, right? Some completion could happen even before we’re finished registering or any time thereafter. We don’t know when this alt is going to complete, when any one of these things we’ve requested are going to finish.

As I said before, making that, at most, one operation completes sound means making the completion of any operation and the disabling of all the other ones in the same alt atomic, that’s the critical thing that has to be made atomic. The other thing that may not be obvious from this, from the description of what alt does, is the fact that there’s going to need to be some cleanup associated with this, and I’ll show you that in a second.

So this takes us to handlers. Essentially, handlers are just like a wrapper around a callback. So, essentially, we’re trying to support real threads where we can block and lightweight go threads and JavaScript where

Handlers

- Wrapper around a callback
- SPI
 - active?
 - commit -> callback-fn
 - lock-id -> unique-id
 - `java.util.concurrent.locks.Lock`: lock, unlock

Figure 15: 00.21.41 Handlers

we can't block or we don't want to block, and we can't block. We don't even have threads. So the least common denominator in all three of these environments in which you want to execute are callbacks because we all know about callbacks. What do we know about callbacks? They are awesome, right?

[Audience laughter]

We love programming directly with callback. No, callback is something you want to hide. But it's okay as an implementation detail, again, of something like this. So handlers are wrappers around callbacks, but they have their own service provider interface, and it's used by the channel to implement the operation so that it all can work.

So the first thing a handler has is this notion of being active. As we said before, we had these handlers to register. We want one to complete, which means the other ones have to be able to say I am no good anymore. Don't use me. And so the active flag is a way to say I'm no longer viable to commit.

We have this commit function, and the commit function takes a handler and just returns the callback. It also, simultaneously, does whatever is necessary to make sure that, if it's one of a set of operations, it's the only one that says yes to commit. So it's just a way to get the callback function out.

And then there are lock IDs associated with every handler, and I'll explain why they're there. And then handlers implement the `java.util.concurrent.lock` API, so a handler is something that you can lock. It's something that the channel can lock.

take/put handlers

- simple wrapper on callback
- lock is no-op
- lock-id is 0
- active? always true
- commit -> the callback

Figure 16: 00.23.39 take/put handlers

So there are very simple handlers associated with take and put. If you look at take, what does take say? Take says I have this function I'd like you to call back when the operation has completed. In the case of the

actual blocking takes, that callback is just going to fulfill a promise and allow the initial thread to proceed. In the case of the go macro threads, that callback is just going to callback through to the state machine or the go macro and allow it to proceed.

So there's nothing else to really do. A take or a put by itself is just its own operation. It's not part of a set of operations, so it's just a wrapper and a callback. There's no lock. The lock operation is a no op. The lock ID is always zero. It's not part of a composite.

Is it active? Is always true. Yeah. I'm ready to go. And commit always just returns that callback, so it's a very small shim used for take and put. Really, all of this SPI is about alt.

alt handlers

- each op handler wraps its own callback, but delegates rest to shared 'flag' handler
- flag handler has lock
 - a boolean active? flag that starts true and makes one-time atomic transition
- commit transitions shared flag and returns callback
 - must be called under lock

Figure 17: 00.24.49 alt handlers

So alt is trickier, right? What we have is a pair of things for every operation inside the alt. There's the operation handler, right? So let's say we'll go back to the scenario where I want a put on A, a put on B, and read from C and take from C. Each of those is an op, is an operation that's going to be part of the alt. They're all going to have their own handler that implements that service provider interface that I talked about before. But what they're going to do is they're going to delegate the locking bits and their commit stuff to a shared flag handler, if you will.

The flag handler has a lock. It has a boolean flag that says this entire alt is still viable to work or the entire alt is now dead, is finished, is completed. It's no longer active. That flag starts true and it makes a one-time transition to false depending on one of the operations committing.

And then committing, any of the operations committing transitions that shared flag to inactive and returns the callback, so that's the atomic commit that disables all of the other things because they're all sharing

this. They all share the same notion of being active, and they're all sharing the same mutex, right? And they're all sharing the same transition. And that's how you can get the multiple operations to do that. But it must be called under the lock of the overall handler, which is again shared.

alt handlers




Figure 18: 00.26.26 alt handlers - build slide

So just to show you that, so here are the handlers, right? So there'll be one of these for each thing in the alt: write, put, put, take. Three purple alt handlers, each of which holds their own callback, and each of which points to a shared implementation of the rest of the handler, which has a single mutex, a single lock ID, and the single flag.

So now, well, if you handed out the first one, and before you even got to the second, handing out, you know, to registering the second one, the first one completed. The second one has no danger of running, right, because the thing it's pointing to has already been toggled to inactive. And that's how we become tolerant of the registration not being atomic.

So the beautiful thing about doing it this way is you don't end up with any global locks, right? We can register everybody independently. You don't have any channel wide multi-locking, right? So when I registered the first operation with channel A, I did not need to have that wait at all for me to do something with channel B. That's completely - it was in A. It was in and out of A's mutex. It just put that operation in the queue, likely, and return.

It's the channel itself that locks multiple handlers because we saw there are these two rendezvous operations, right? The two rendezvous, right? In the first put diagram, we saw the one thing that launched commit. And in the take diagram, we show the other thing that launched commit. And it's those two operations that actually involve something being atomic on two handlers.

alt concurrency

- no global or multi-channel locking
- but channel does multi-handler locking
some ops commit both a put and take
- lock-ids used to ensure consistent lock acquisition order

Figure 19: 00.27.14 alt concurrency

I want to tell the taker you got it; your take was accepted. And I want to tell the putter, your put was accepted and your put went to this take. And both of those things, both of those handlers need to be told atomically that the transfer has happened. And so what happens is both of the handlers get locked, and then we can commit and commit, and then take the locks off the handlers. So there are some operations that do both of them, and that's the explanation for the lock IDs.

I have random sets of operations: A, B, C here; A, B, D there; D, B, C, A. Okay, so what order should I lock these things? If I'm using B and A, and you're using A and B, that's a recipe for what?

[Audience member: Deadlock.]

Deadlock, right? So the classic technique for dealing with deadlock when you only really have non-nested locking, when you just have a locking set, have a set of things to lock, is to use consistent lock acquisition order. In other words, no matter what the set is, we're going to order them consistently. Then we can lock in the same order, and we can't deadlock that way.

And that's what happens, right? The lock IDs are assigned monotonically to these handlers. And, during one of those multi-handler operations, their lock IDs will be examined, and they'll be done least and greater no matter what by lock ID. And that means lock acquisition order is consistent and the system can't deadlock. So it's simple, but that's how it's done.

alt cleanup

- 'disabled' handlers will still be in queues
- channel ops purge

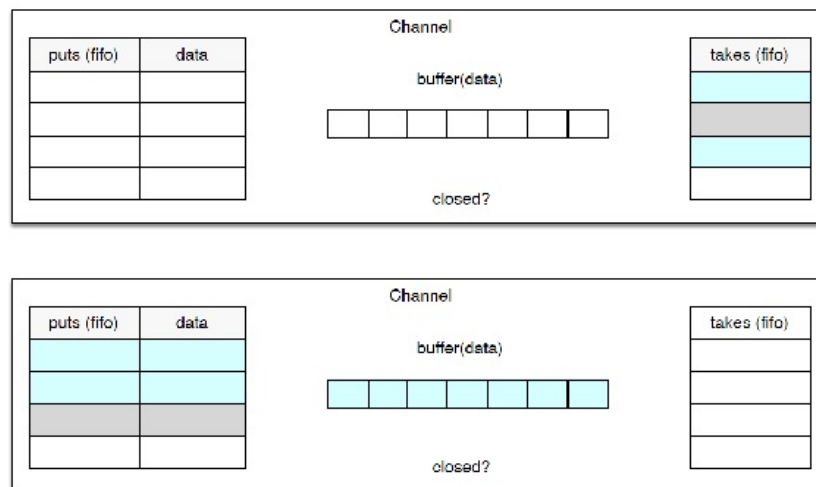


Figure 20: 00.29.44 alt cleanup

So, now we can get more of the truth of these diagrams. So it ends up that - so what happens? What happens if I just did an alt and everything was in queue, nothing succeeded right away, nothing completed

right away? And then one of the alt operations completed. What happens to the other ones? Where are they?

They're somewhere in these queues, right? They have to be in these queues. They got registered, and we're waiting for them to happen, and now they're in these queues. But they're inactive. When it comes time to do them, there'll be nothing to do.

And so being - that active flag is used and it can be called outside of a lock, by the way, is used to determine whether or not something is still a candidate for completion. And if it's not, it needs to get purged. And so what happens is the ordinary channels ops, the ordinary takes and puts will go into these queues and clean up the dead requests, the dead operations. And that's how that's done, so there's nothing for you to do, and they'll just get purged automatically. But they are in those queues until that happens.

SPI revisited

- handler callback only invoked on async completion
 - only 2 scenarios
- when not 'parked', op happens immediately
 - *callback is not used*
 - *non-nil return value is op return*

Figure 21: 00.30.57 SPI revisited

So I talked about this completion of a handler, and now we know better about how handlers work, so we want to look more carefully at the two scenarios when this is happening. It ends up that there are only two scenarios out of the six plus two that I showed where there is a two-element completion under any kind of callback of the handler. All of the other operations, so if you enter the operation and I said - and it immediately completes, the callback is not called in those cases. It's not used.

If you're not parking the operation, that return value of the operation is returned immediately by the operation. Now remember, this is the service provider interface to the wrapping operation. So the wrapping operations, you know, arrow bang (->!), and things like that, they look at the return value and say, "Oh, you know what? You just immediately returned to me. I'm not going to look at that promise because I

know my callback is never going to be used. And the callback was the only thing that ever fulfilled that promise, and I'm just going to ignore that promise and return to the user."

When you get a non-nil value, it means that the operation actually returned. When you get a nil value from the service provider interface, it means your operation was in queue and your callback will be called eventually.

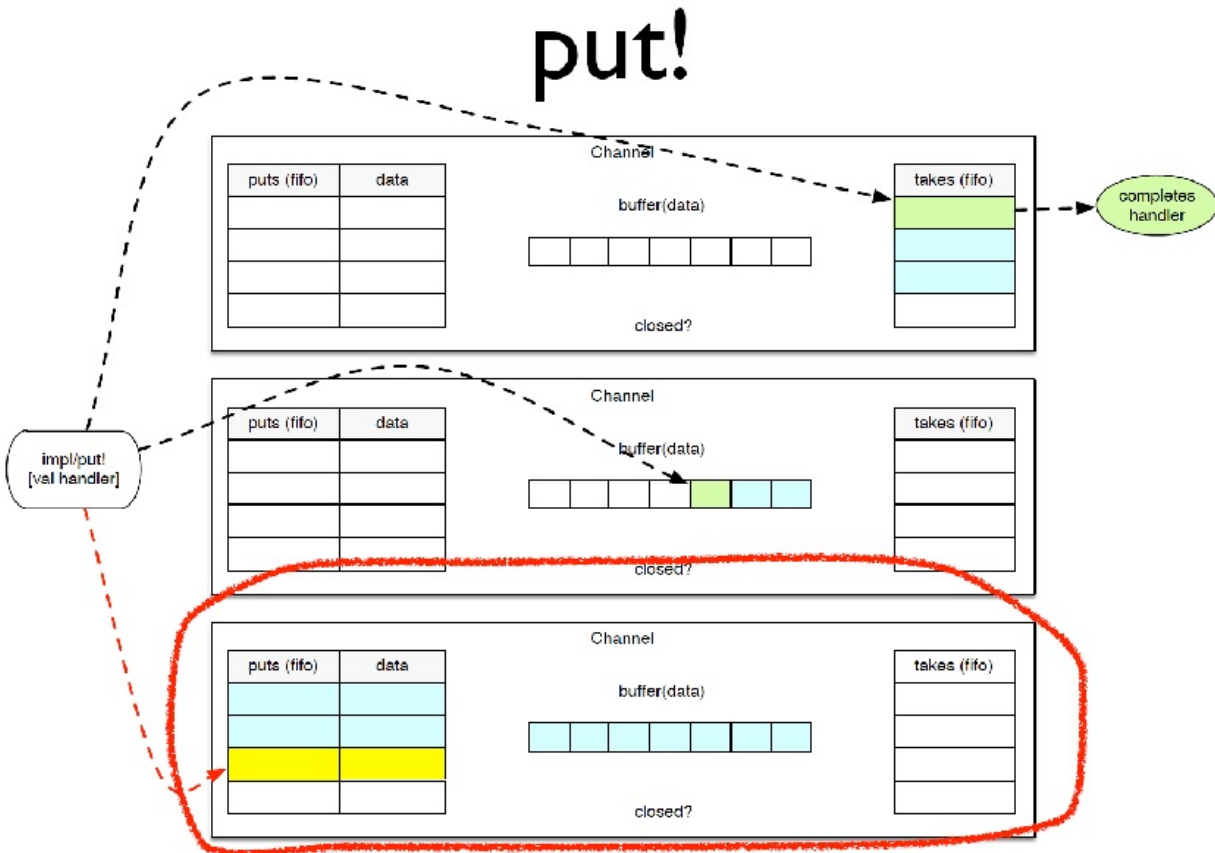


Figure 22: 00.32.20 put!

So, if we go back to look at these diagrams, this is the case for put and the only case for put where this enqueueing is going to go on where it gets blocked, and the same thing there. And, similarly, you have to sort of switch between these, right? The case where you got enqueued for put is the place where you get completed, right? That same bottom scenario is the place where the put gets completed by a take later. Right? So for both the blocking cases, the corresponding completion case is the opposite operation.

So here again, the take that gets blocked is in the first case, which is the one that completes when a put is done. It just then completes the take. So the block and the completions pair up like that. And only in those two cases do you get the handlers running.

All right, so then there's the wiring of this up, right? So that's the service provider interface. That's not what you actually use. Then we have sort of two implementations that fill in the details of what arrow bang (->!) and double arrow or arrow double-bang (->!!) and vice versa do.

In the blocking set, so in the thread macro and the double-bang (!) operations, a promise is always created, so the consumer of the service provider interface, which is the API function, the API function creates a

take!



Figure 23: 00.33.00 take!

Wiring !/!!

- blocking ops (!!)
create promise
callback delivers
only deref promise on nil return from op
- parking go ops (!)
IOC state machine code is callback

Figure 24: 00.33.22 Wiring!/!!

promise. And it makes a callback that completes, that delivers that promise - usually with a return value on a take. For put, it just delivers the promise.

And then it tries the operation, right? If it gets back non-nil, it knows that's the actual return value. A callback is not going to be used, and it will not de-ref that promise. If it gets back nil, then it tries to de-ref the promise, and in which case it's blocked.

For the parking ops, the go ops, what happens is the callback that's registered is the state machine resumption callback. And, similarly - actually, I'm not sure if ... takes advantage of the fact that the direct return is available. He may always wait for the callback, but I don't believe he's allowed to do that, so that's how it gets wired up.

[Audience laughter]

Summary

- You don't need to know any of this
- But understanding the 'machine' can help you make good decisions



Figure 25: 00.34.55 Summary

You don't need to know any of this, but understanding it, I think, can help, and I want to leave enough time for questions, which it seems like I have, so any questions? Yes.

[Audience member: One of your best nights, you talked about the puts, I mean, I don't know ... use hooks ... multiple channels. Can you describe the scenario or expand a bit more why you use...?]

No, alt can do puts or takes.

[Audience member: Yeah ... by why? What's the rationale?]

Versus?

[Audience member: I can understand when you ... take....]

Yeah.

[Audience member: It's like ... select. You want to know kind of where ... came from.]

Yes.

[Audience member: I'm just trying to get my head around when you use or when you might use put.]

Oh, okay. So when would you use alt around puts?

[Audience member: Exactly.]

When you have consumers with different capabilities, right? So maybe you have a set of consumer channels I can communicate with, and one of them is slower than the others. So I want to try to write to everybody, but whoever is read will be the first one to complete. So maybe I'm issuing operations that take varying amounts of time. So I have four people that are going to help me out, do a bunch of jobs that take varying amounts of time, and I now have a new job, who should I ask?

I don't know. There's not really a good answer to that question, but alt is a good answer to that question. I'm going to ask the first one. Who can say yes? That's what alt lets you do. And that's a lot of power, so that's why. That would be one case why. Does that make sense?

[Audience member: Yes.]

Okay.

[Audience member: So what's the difference with all these workers consuming the same queue where ... put in on the same channel? Because reverse that, all these four consumers ... taking on one channel...?]

There could be, but the beautiful thing is that, for that simplest scenario, that might work okay.

[Audience member: Yeah.]

There may be other semantics associated with the operations. You have priority to combine with it. You may have a cost metric associated with different things. You may want to try one of these three and then a fourth. Also, the way you compose them, it may not have been available to you to make them all consumers of the same channel. They all gave you channels, right? So you have sort of a bunch of options for putting things together. Sometimes you can achieve the same result two different ways.

Other questions? Yes.

[Audience member: There's something like ... why does...?]

Well, it would be easy. That's for sure. The question was why wouldn't a global mutex be good enough. It would be great. I mean a global mutex could always solve any problem.

[Audience laughter]

As long as all of your program is ready to wait for all of the rest of your program.

[Audience laughter]

The thing is that there are multiple problems associated with this. So a global mutex could make registration atomic, right? But that doesn't help you because there are two other hard jobs you have to do. One is you have to make disabling all the other alts atomic also, and you have to make rendezvous atomic. So unless you want to keep going back to that global mutex and do all three of those jobs that way, it may – I don't like it as a solution for any.

STM doesn't do it, and I think you don't want to build concurrent systems that have global mutexes because it makes too many things wait that needn't, right? You can have two completely different sets of channel operations that are unrelated to each other. Why should they be making each other wait? We're dealing with different sets of channels. We shouldn't be interacting with each other at all. So I think global mutexes, in general, are bad. They are easy though, for sure. But it ends up, in this case, they're really a not very

good solution to registration and a much worse solution for the other two things: for atomic completion and for rendezvous. But, yeah, you could do the job that way.

In all cases where you use multiple locks you could always say, if there was one lock this would be easy, and it would be. But it wouldn't allow things to be concurrent that could otherwise be concurrent. And every time anybody has a global lock, everybody hates it eventually. It's just a matter of time because it is in the way, needlessly. So it's ruled out by my esthetic sense.

[Audience laughter]

Not by technical reasons, let's say.

[Audience member: Hello.]

Hi.

[Audience member: I was thinking about....]

Yes

[Audience member: (Indiscernible)]

No.

[Audience member: ...10,000 or 100,000....]

[Audience member: Yeah ... I just wanted to know is it outcome or...?]

No, no. It shouldn't be that.

[Audience member: It shouldn't be that. Okay.]

No, he was probably using a sliding buffer or something.

[Audience member: (Indiscernible)]

Yeah. I'm not sure exactly what –

Other questions? Yes.

[Audience member: Do you think that the buffer and queue sizes are useful metrics to monitor, or is that more like...?]

No, they would be great to monitor, and it's definitely on the to-do list to add monitor ability of those things. Yeah, you would. It's our intent to make the channels more evident so that you could look at them, so you can make decisions and get warnings and things like that. That's definitely an enhancement area. Yes?

[Audience member: ...core.async?]

That's hard to say. Certainly the monitoring would be great. I think there are some more nuance notions of channel, buffer policies that are useful. You know, having a pluggable policy, I think, is just generally useful because there are all kinds of logic you could have about the context of the buffer. These messages are higher priority messages than those would be a great example. So I think that that's an active area or an area that's ripe for enhancement: more sophisticated buffers. And that's maybe the first part of the service provider interface that we try to finalize for exposure because I don't think we're going to provide every kind of buffer policy you could want, and it's probably a great area for you to customize for yourself.

Another thing that's interesting to me, and I can't promise I'll do it or I've started working on it other than the early design stuff is, I think core.async has already proven its utility to many of us. Immediately people started using it and having success with it, and it's just working. So it's become important, I think, to me, the stuff I write, and to a lot of people who are doing production work. And I think that the go macro, in all of its glory, is a great proof of concept of what you could do with a macro with now several thousand lines of code behind it, that whole mini sub-compiler and whatever. But it's quite challenging to get that right. I think there are still a lot of things in there that are difficult because it's trying to sort of re-implement Clojure.

It also does the translation of Clojure to Clojure to solve the problem of doing the state machine. And some of the raw materials you might want to use to do that are lower level than Clojure is itself, like go-to byte code would be useful for it, and just lower level access to the way the compiler admits local variable access. So one potential solution or enhancement to that would be to actually build async support into the compiler and make not really go blocks. Go blocks would be a level higher, but the meat and potatoes of a go block is the set of code that interprets Clojure code in a way that doesn't use local variables on the stack, right?

The idea of all that's happening in the go macro is that all the things that were local variables become data. You can imagine, you could turn all the local variables and move them from the stack to fields of the method objects that you saw before. When you've done that, you now have the ability to say, well, I have nothing on the stack. I could just – I could stop now if I wanted to, and I could tell somebody to resume me if I had a way to remember where I am and, when I'm next called, jump or go to that spot.

And so it would be relatively straightforward, I think, to enhance the compiler to compile what I'll call async blocks that would implement local variables via fields and would have a generic way for you to declare a function as being async, which means it's going to comply with this service provider interface the way I'm handling these callbacks to say I'm going to hand you a callback. If I can immediately consume, I'm going to return a result. If I can't, I'm going to return nil, and I will definitely call your callback later with the result. And with that contract in place, you could build the go macro and its support for the IO ops of Clojure, of core.async, as it stands now, but you could also build things like generators within yield and some other generic thing. Having it be part of the compiler proper would mean it's the same implementation of all the code and likely to be, you know, to work right away. And it would just mean some conditionality on the code that admits local variable access.

So I think it's become important enough that the pride moment of, look, you can do this with just a macro, is not dominated by the production objectives of having something that's robust and as fast as we can make it, which I think this would be. So that's a candidate enhancement. That's a very long answer to that question. I wanted to explain it because I'm sure Tim will watch this and be like, "Oh, my God. My macro!"

[Audience laughter]

[Audience member: What you were just describing sounds a bit like....]

No, no. Uh-uh.

[Audience laughter]

[Audience member: (Indiscernible)]

Well, so a continuation is definitely more general. And this is not going to use continuation passing style or anything like that, so I'm loathed to call it a continuation. It's in the same category, certainly. I'm joking with you, but I think it's best to reserve the term continuation of something that really is a continuation, which this would not be. And the operation provided by async would not be like call CC in that way in terms of the continuation being first class. It wouldn't be – you couldn't resume it more than once and a bunch of other limitations will be present, so it's really for a very specific set of use cases, which are very general. I mean, I just say Ollie give a talk about just generators are enough to do what a lot of people think they need a lot more stuff to do.

Yeah, it's in the same category, but I think continuations proper are pretty difficult for people to understand and difficult to use. And implementing with our full set of implications is also a challenge that can compromise other aspects of your real time performance, so Clojure is not going to get continuations.

[Audience laughter]

[Audience member: Was there something planned for ... binding ... core.async? There are a bunch of issues using those....]

With go?

[Audience member: With dynamic binding and....]

With the go macro? Specifically, yeah, so – yeah, I mean there are functions that allow you to do the conveyance. I’m not sure the go macro can allow all those to work yet. But certainly the thing I was talking about wouldn’t have to be any different. It would still be at a level higher up to deal with that conveyance of bindings.

Yes?

[Audience member: If I remember correctly, Timothy Baldridge released recently something with channels on the network. Could you tell us something more about that?]

Hmm.

[Audience laughter]

So it’s easy to have something that you call a channel and put over a wire. But it’s actually pretty hard to get all the semantics of these channels to work over wires. And there are lots of wires that already exist. You already have queues. You already have all kinds of interfaces for doing stuff over wire, and so I’m not sure that we need to say this is a channel over wire. And I especially don’t want to say this is a channel over wire if it’s not actually the semantics of channels. And so things like atomic, one time only completion of alt over more than one wire, I don’t think so.

So it’s possible, for instance, to have a subset of channels that work over wires, but you have to know what they are and you have to be able to talk about them. If you look at the inside of channels and the implementation, you’ll actually see this word called port. And there’s a semantic notion, which is not exposed anywhere other than in the naming of ports and sort of having an endpoint, a read endpoint or a write endpoint independent of each other. And I think we could come up with semantics for ports where we might be able to say that you could have the read semantic of a port over a wire, or we could have limitations on alt.

But I really – and I discourage Tim from doing that and still do. But he understands these are my reservations about it. I think that there are a lot of people – how many people are using core.async and having a wire somewhere in the middle? It’s okay to admit. I am.

[Audience laughter]

But the thing there is to understand the wire; it has its own semantics. It’s got an API. Maybe it uses callbacks. Maybe it has some semantics about failure, message delays, queuing, and everything else. There’s no crime in saying, okay, well, something is going to come over a wire, and that will have the semantics it has. And I will take whatever it produces with whatever semantics it has, and I will put that thing on a channel.

And when I’ve done that, now I have two things I can understand. I can go look at the documentation of the wire and say this wire is not working. And what’s wrong? And I’m using some ZeroMQ or HornetQ or HTP or some other thing, raw sockets. This isn’t working. What is it supposed to do? I understand that. It did come over. Now it ended up in the channel. Now something went wrong. Now I have semantics of channels. That all makes sense to me. That’s something I can build and compose on.

And I don’t think there’s anything wrong with that. I think that’s actually better than pretending that you had these semantics over wires, when in any way you don’t. But that’s not to say that we won’t end up with a subset of the things that work over wires, but it’s so easy to take something that came over a wire and call it put bang (put!), and in fact that’s really what I recommend you do. You’re going to get an IO library, right? It’s going to have a callback. That’s a perfect place to put bang (put!) and enter the channel world and just be channels from now on.

I like that separation myself. That’s what we’re doing in production code that uses channels is we’re not pretending the channel goes over the wire. You can write to a channel and say that channel will eventually talk to a wire, but the semantics of the wire are what they are.

Yeah?

[Audience member: Is there a typical way to monitor a go...?]

No.

[Audience laughter]

[Audience member: ...some time?]

What kind of monitoring do you want to do?

[Audience member: See that it's still working, still alive, maybe kill it if it's not working.]

Yeah, so I mean I think that we talked already about monitoring of the channels themselves, so I think to the extent a go block is dependent on the channels its interoperating with, that gives you some insight. The killing part, well, you can't really make an operation that would kill channels and threads, equivalently, because you can't kill threads in the same way, but you do have more ability, for instance, to tell a channel to stop working.

I'm not sure I would want an out of band extra thing that did that. People are building go blocks and processes that live inside a go block that are stoppable. And the nice, super clean way to do that is to have a channel on which you communicate stop. And I think, in almost all cases, that's better. It's not really the same thing as monitoring, but you'd have to tease out what it is you want to look at. Killing it, I think, is something better done with the channel itself. But if the channels were monitorable, then you could see how is it successfully producing and consuming the way you want by looking at the channels themselves.

[Audience member: Yeah, basically like, you know, thinking that you can monitor your system and—]

Yeah, if there's arbitrary code in there, I mean there's arbitrary code in the rest of your systems too, which you can't monitor necessarily either, right?

[Audience member: Yes.]

[Audience member: Last one.]

Last one. Yes?

[Audience member: What other options you have that might be during the design and what were the other options you rejected ... implement core.async?]

Implementation options?

[Audience member: ...the design, like other options than the...]

Something other than CSP or—?

[Audience member: Yeah.]

At that level you're saying options? Yeah, so I've seen the generator stuff. Obviously there are continuations. I really like what the go language did. I thought they made a good choice and, you know, they're building on top of what people have been doing with CSP all along. It was the Java CSP library that implements the same kinds of operations and it didn't seem to me to be a great area to invent because it's difficult to get the semantics of these things correct and really have stuff that's useful as a building block. And I didn't see anything fundamentally wrong. I mean, the thing I wanted to do was make sure alt was a regular function and not syntax. You know, so there's that alt function, which is quite powerful for that to be a function. That's an enhancement versus go, but there are small things like that.

I think it's just a fabulous fit for Clojure because, in Clojure, we not only have this, but we combine it with the fact that what we're putting on these channels is immutable. And so that combination of things really yields something that's robust. And that's why I think it's just a super great, super great for Clojure, and I'm very happy to see everybody using it.

Given that was the last question, I just want to use my opportunity up here to say I am so happy to see everybody here and see Clojure grow and to see what everybody is doing with it. It's really fantastic, and it makes me very happy, so congratulations to all of you for what you're accomplishing with Clojure.

[Audience applause]